




DATABRICKS

Unified Data Access with Spark SQL

Michael Armbrust – Spark Summit 2014

@michaelarmbrust

Spark SQL Components



- Catalyst Optimizer
 - Relational algebra + expressions
 - Query optimization
- Spark SQL Core
 - Execution of queries as RDDs
 - Reading in Parquet, JSON ...
- Hive Support
 - HQL, MetaStore, SerDes, UDFs

Relationship to

Shark modified the Hive backend to run over Spark, but had two challenges:

- » Limited integration with Spark programs
- » Hive optimizer not designed for Spark

Spark SQL reuses the best parts of Shark:

Borrows

- Hive data loading
- In-memory column store

Adds

- RDD-aware optimizer
- Rich language interfaces

Migration from

Ending active development of Shark

Path forward for current users:

- Spark SQL to support CLI and JDBC/ODBC
- Preview release compatible with 1.0
- Full version to be included in 1.1

<https://github.com/apache/spark/tree/branch-1.0-jdbc>

Migration from

To start the JDBC server, run the following in the Spark directory:

```
./sbin/start-thriftserver.sh
```

The default port the server listens on is 10000. Now you can use beeline to test the Thrift JDBC server:

```
./bin/beeline
```

Connect to the JDBC server in beeline with:


```
beeline> !connect jdbc:hive2://localhost:10000
```

*Requires: <https://github.com/apache/spark/tree/branch-1.0-jdbc>

Adding Schema to RDDs

Spark + RDDs

Functional transformations on partitioned collections of **opaque** objects.



SQL + SchemaRDDs

Declarative transformations on partitioned collections of **tuples**.

Name	Age	Height
Name	Age	Height
Name	Age	Height
Name	Age	Height
Name	Age	Height
Name	Age	Height

Unified Data Abstraction



Using Spark SQL

SQLContext

- Entry point for all SQL functionality
- Wraps/extends existing spark context

```
from pyspark.sql import SQLContext  
sqlCtx = SQLContext(sc)
```


Example Dataset

A text file filled with people's names and ages:

```
Michael, 30
```

```
Andy, 31
```

```
Justin Bieber, 19
```

```
...
```

RDDs into Relations (Python)

Load a text file and convert each line to a dictionary.

```
lines = sc.textFile("examples/.../people.txt")
```

```
parts = lines.map(lambda l: l.split(","))
```

```
people = parts.map(lambda p: {"name": p[0], "age": int(p[1])})
```

Infer the schema, and register the SchemaRDD as a table

```
peopleTable = sqlCtx.inferSchema(people)
```

```
peopleTable.registerAsTable("people")
```

RDDs into Relations (Scala)

```
val sqlContext = new org.apache.spark.sql.SQLContext(sc)
import sqlContext._

// Define the schema using a case class.
case class Person(name: String, age: Int)

// Create an RDD of Person objects and register it as a table.
val people =
  sc.textFile("examples/src/main/resources/people.txt")
    .map(_.split(","))
    .map(p => Person(p(0), p(1).trim.toInt))

people.registerAsTable("people")
```

RDDs into Relations (Java)

```
public class Person implements Serializable {  
    private String _name;  
    private int _age;  
    public String getName() { return _name; }  
    public void setName(String name) { _name = name; }  
    public int getAge() { return _age; }  
    public void setAge(int age) { _age = age; }  
}
```

```
JavaSQLContext ctx = new org.apache.spark.sql.api.java.JavaSQLContext(sc)  
JavaRDD<Person> people = ctx.textFile("examples/src/main/resources/  
people.txt").map(  
    new Function<String, Person>() {  
        public Person call(String line) throws Exception {  
            String[] parts = line.split(",");  
            Person person = new Person();  
            person.setName(parts[0]);  
            person.setAge(Integer.parseInt(parts[1].trim()));  
            return person;  
        }  
    });
```

```
JavaSchemaRDD schemaPeople = sqlCtx.applySchema(people, Person.class);
```

Querying Using SQL

*# SQL can be run over SchemaRDDs that have been registered
as a table.*

```
teenagers = sqlCtx.sql("""  
    SELECT name FROM people WHERE age >= 13 AND age <= 19""")
```

*# The results of SQL queries are RDDs and support all the normal
RDD operations.*

```
teenNames = teenagers.map(lambda p: "Name: " + p.name)
```

Caching Tables In-Memory

Spark SQL can cache tables using an in-memory columnar format:

- Scan only required columns
- Fewer allocated objects (less GC)
- Automatically selects best compression

```
cacheTable("people")
```

Language Integrated UDFs

```
registerFunction("countMatches",
```

```
    lambda (pattern, text):
```

```
        re.subn(pattern, '', text)[1])
```

```
sql("SELECT countMatches('a', text)...")
```

SQL and Machine Learning

```
training_data_table = sql("""
    SELECT e.action, u.age, u.latitude, u.longitude
    FROM Users u
    JOIN Events e ON u.userId = e.userId""")

def featurize(u):
    LabeledPoint(u.action, [u.age, u.latitude, u.longitude])


// SQL results are RDDs so can be used directly in MLLib.
training_data = training_data_table.map(featurize)

model = new LogisticRegressionWithSGD.train(training_data)
```


Hive Compatibility

Interfaces to access data and code in the Hive ecosystem:

- Support for writing queries in HQL
- Catalog info from Hive MetaStore
- Tablescan operator that uses Hive SerDes
- Wrappers for Hive UDFs, UDAFs, UDTFs



Reading Data Stored in Hive

```
from pyspark.sql import HiveContext  
hiveCtx = HiveContext(sc)
```

```
hiveCtx.hql("""  
    CREATE TABLE IF NOT EXISTS src (key INT, value STRING)""")
```

```
hiveCtx.hql("""  
    LOAD DATA LOCAL INPATH 'examples/.../kv1.txt' INTO TABLE src""")
```


Queries can be expressed in HiveQL.

```
results = hiveCtx.hql("FROM src SELECT key, value").collect()
```

Parquet Compatibility

Native support for reading data in Parquet:

- Columnar storage avoids reading unneeded data.
- RDDs can be written to parquet files, preserving the schema.



Using Parquet

*# SchemaRDDs can be saved as Parquet files, maintaining the
schema information.*

```
peopleTable.saveAsParquetFile("people.parquet")
```

*# Read in the Parquet file created above. Parquet files are
self-describing so the schema is preserved. The result of
loading a parquet file is also a SchemaRDD.*

```
parquetFile = sqlCtx.parquetFile("people.parquet")
```

Parquet files can be registered as tables used in SQL.

```
parquetFile.registerAsTable("parquetFile")
```

```
teenagers = sqlCtx.sql("""  
    SELECT name FROM parquetFile WHERE age >= 13 AND age <= 19""")
```

Features Slated for 1.1

- Code generation
- Language integrated UDFs
- Auto-selection of Broadcast (map-side) Join
- JSON and nested parquet support
- Many other performance / stability improvements

Preview: TPC-DS Results

